

Blocked Dynamic Sparse Attention

Mehul Goel (mehulg), Saransh Agrawal (saransh2)
Website: <https://mehulgoel873.github.io/15418-final-project/>

Summary

Background

Attention




Figure 1: The attention block. Three projections (MatMuls (1)–(3)) produce Q , K , V ; MatMul (4) and the softmax operate on an $N \times N$ matrix (highlighted), and MatMul (5) collapses back to $N \times d$.

Modern transformers process a sequence of tokens by repeatedly applying an *attention* operation. Each token is represented as a d -dimensional embedding vector, so the input to an attention layer is a matrix $\mathbf{x} \in \mathbb{R}^{N \times d}$, where N is the number of tokens (the *context length*) and d is the embedding dimension.

Figure 1 sketches the full attention block. The input $\mathbf{x}$ is first projected by three learned $d \times d$ weight matrices $\mathbf{W}_q$, $\mathbf{W}_k$, $\mathbf{W}_v$ (MatMuls (1)–(3)) into a query Q , a key K , and a value V , each of shape $N \times d$. The layer then computes

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right) V.$$

MatMul (4) forms the $N \times N$ score matrix QK^T , a row-wise softmax turns these scores into normalized weights, and MatMul (5) multiplies them by V .

The two highlighted stages in Figure 1 scale with the *square* of the sequence length. MatMul (4) materializes an $N \times N$ matrix, the softmax sweeps over it row by row, and MatMul (5) reads it

again. In practical settings $N \gg d$. Today’s frontier models target context lengths of a million tokens or more, while d is typically only 64 or 128 per head, so a single attention layer would touch on the order of 10^{12} entries in its intermediate matrix. Both the arithmetic and the memory traffic of storing and revisiting an $N \times N$ matrix become impractical at this scale.

Each block in the figure maps well to a GPU on its own. The matrix multiplications are embarrassingly parallel: each output entry can be computed independently. The softmax requires a reduction along each row, but the rows themselves are independent and parallelize across N . The bottleneck is therefore not a lack of parallelism but the volume of computation and data movement on an $N \times N$ matrix at long context lengths. Most entries of the softmax output are close to zero (highlighted as “sparse” in Figure 1), and that observation motivates the tiled and sparse approaches we explore in the rest of this report.

Sparsity

A long line of work has shown that the score matrix $\sigma(QK^T)$ produced by attention is approximately sparse: most of its entries are very small after the row-wise softmax, and the output of MatMul (5) is dominated by a small fraction of the columns. This has motivated a broad family of *sparse attention* schemes that fall into two camps. Static patterns (such as local windows, strided, or block-diagonal masks) fix a sparsity structure ahead of time and are easy to map onto regular GPU tile layouts. Dynamic patterns (such as top- k or learned routing) choose which entries to keep based on the input itself, often giving better accuracy at a given sparsity level at the cost of more irregular computation [1].

Rather than commit to a single pattern, our experiments use a randomly-generated sparse mask. Each cell is independently kept with probability p , where p is the fraction of dense entries. This lets us sweep sparsity levels without entangling our results with the biases of any particular method. To match the granularity at which real-world transformers exploit sparsity, the mask is defined at a *block* level rather than per-element: a granularity of g means we partition the $N \times N$ score matrix into $g \times g$ tiles and label each tile as fully dense or fully sparse. A granularity of 1 corresponds to per-element sparsity, while a granularity of 32 means 32×32 tiles are the smallest unit of sparsity.

Block-level sparsity reduces the total amount of work on the score matrix, but it gives up structure that dense kernels rely on. Coalesced memory accesses, vectorized loads, and predictable load balancing all assume regular tile traversal; once tiles are skipped according to a mask, warps within a thread block can be assigned very different amounts of work, the rows of the score matrix lose alignment in memory, and arithmetic intensity drops. Recovering performance therefore requires both a sparse-friendly storage format and kernels that revisit the three stages of attention with the mask in mind.

BCSR storage. We store the block-level mask in *Blocked Compressed Sparse Row* (BCSR) format. BCSR is the block analogue of CSR: instead of recording the position of every individual nonzero, it records the position of every nonzero *tile*. For each row of tiles, we store the column indices of the dense tiles, and the dense tiles themselves are laid out contiguously in memory in row-major order. This keeps the inner loop of every kernel operating on dense, contiguous tiles (so coalescing and vectorization still apply within a tile) while cutting work in proportion to the fraction of sparse tiles.

Sparsity in MatMul. The two matrix multiplications in Figure 1 change shape under block sparsity. MatMul (4) computes QK^T , but only at positions where the mask is dense. This is the *sampled dense-dense matrix multiplication* (SDDMM) pattern: Q and K are dense, and the

mask selects which output tiles to compute. Skipping a sparse output tile skips an entire block of dot-product reductions, which is where most of the savings come from. MatMul (5) computes $\sigma(QK^T) \cdot V$, where the left operand is now a block-sparse matrix in BCSR form and V is dense. This is the *sparse-dense matrix multiplication* (SpMM) pattern: each dense tile of the score matrix contributes a small dense matmul with the corresponding rows of V , and the BCSR row indexing tells each thread block which tiles to accumulate.

Sparsity in Softmax. The row-wise softmax sits between MatMul (4) and MatMul (5), and it has to respect the same mask. Sparse tiles can be treated as containing $-\infty$ scores, so they contribute zero to both the per-row sum and the output. We never materialize those tiles in practice: the softmax kernel walks the BCSR row directly, computes the per-row max and the exponential sum across only the dense tiles in that row, and writes back the normalized values in place. The reduction structure is the same as in dense softmax, but the length of each reduction is now proportional to the number of dense tiles in the row rather than to N .

Approach

Results

References

- [1] Tianyang Lin, Yuxin Wang, Xiangyang Liu, and Xipeng Qiu. A survey of transformers, 2021.

List of work

The total credit for the project was 50-50 between Saransh and Mehul.